

The background of the slide is a dark navy blue, filled with a bokeh effect of out-of-focus light circles. These circles vary in size and are primarily a warm yellow-gold color, with some cooler blue and teal tones interspersed, particularly towards the top and bottom edges. The overall effect is a soft, textured, and modern aesthetic.

DEEP LEARNING FRAMEWORKS

Contents



- Introduction to Deep Learning Frameworks
- Tensorflow & Keras
- Pytorch
- Other frameworks & Interoperability

Introduction

Deep Learning Frameworks

- Set of tools to train neural networks
- Optimized code
- Parallel processing (CPUs, GPUs, TPUs)
- Community support (new architectures, datasets, ...)

Choosing a Framework

- Today, most of them have the same capabilities.
- Ease of use (understanding what you are doing)
- Ecosystem (research, development, production)
- Target application (web, server, IoT, ...)
- Accessibility and flexibility for professionals from other fields.

The background of the image consists of numerous overlapping circles in various shades of teal and blue. The circles vary in size and opacity, creating a bokeh-like effect. The colors range from a deep, dark blue to a lighter, vibrant teal. The overall composition is abstract and modern.

Tensorflow



- Most used framework in the industry
- Developed by Google
- High-level API (Keras)
- Low-level API (TF2.0)
- Great production support (web, mobile, servers, IoT, ...)

Keras

```
# import tensorflow and keras
```










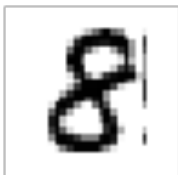


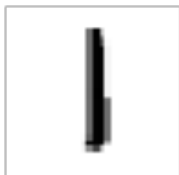
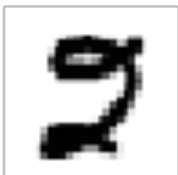


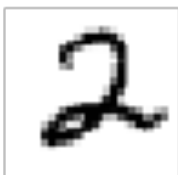





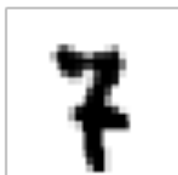








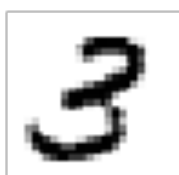
```
import tensorflow as tf  
from tensorflow import keras
```

```
# download a dataset
```

```
mnist = keras.datasets.mnist  
(X_train_full, y_train_full), (X_test, y_test) = mnist.load_data()
```

```
# split and normalize
```

```
X_valid, X_train = X_train_full[:5000] / 255., X_train_full[5000:] / 255.  
y_valid, y_train = y_train_full[:5000], y_train_full[5000:]  
X_test = X_test / 255.
```


7 	3 	4 	6 	1 	8 	1 	0 
9 	8 	0 	3 	1 	2 	7 	0 
2 	9 	6 	0 	1 	6 	7 	1 
9 	7 	6 	5 	5 	8 	8 	3 

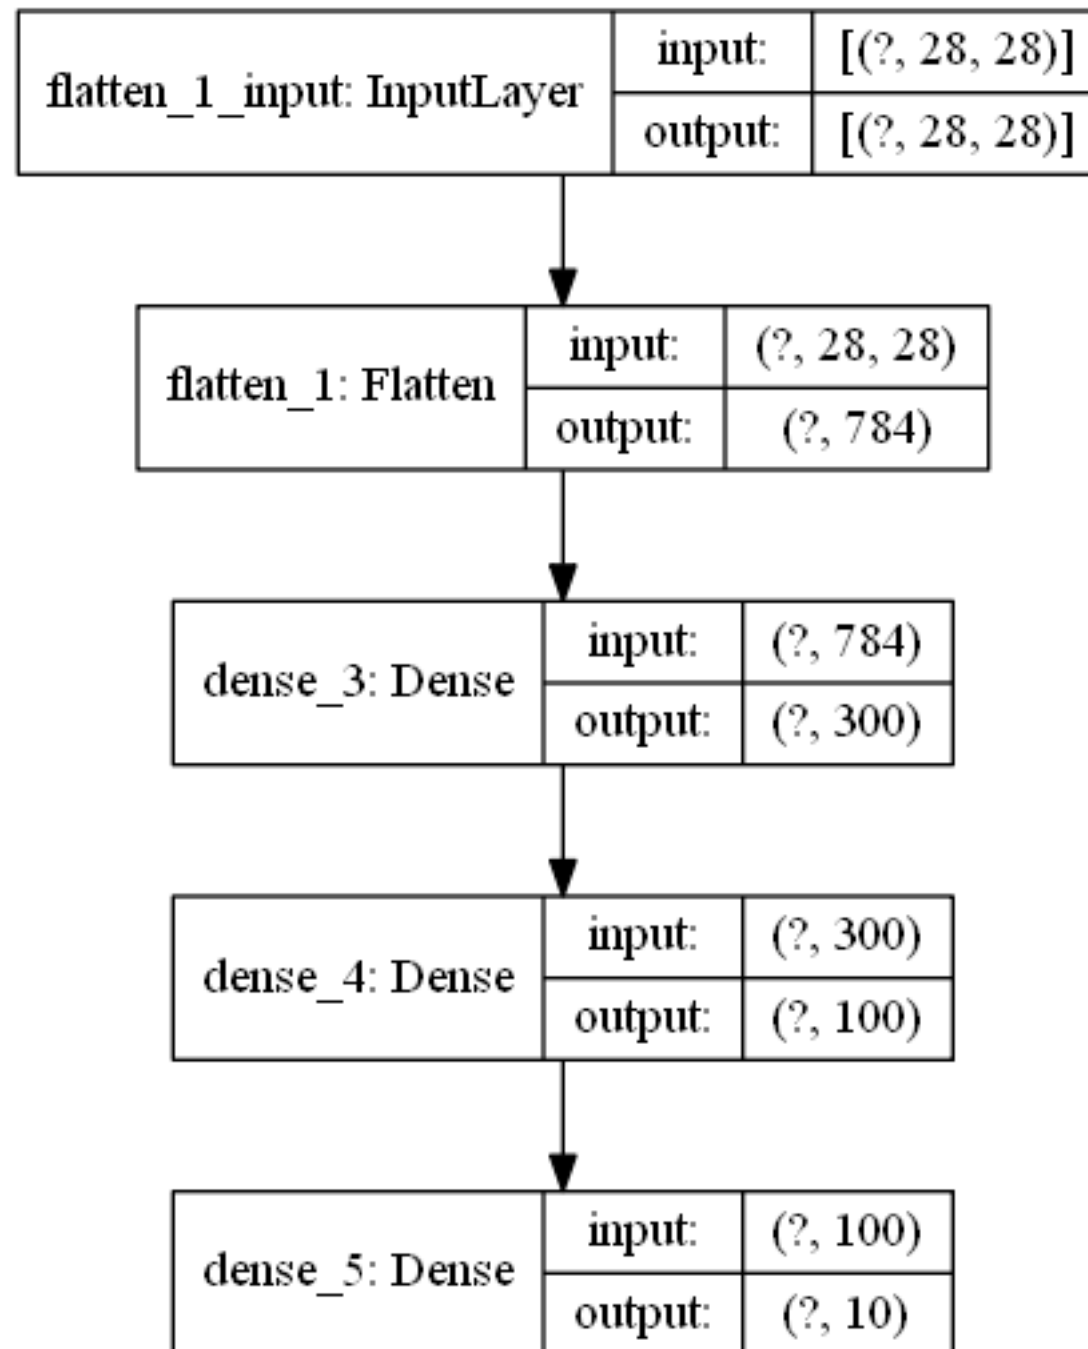
```
# build a model
```

```
model = keras.models.Sequential([  
    keras.layers.Flatten(input_shape=[28, 28]),  
    keras.layers.Dense(300, activation="relu"),  
    keras.layers.Dense(100, activation="relu"),  
    keras.layers.Dense(10, activation="softmax")  
])
```

```
model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====	=====	=====
flatten_1 (Flatten)	(None, 784)	0
dense_3 (Dense)	(None, 300)	235500
dense_4 (Dense)	(None, 100)	30100
dense_5 (Dense)	(None, 10)	1010
=====	=====	=====
Total params: 266,610		
Trainable params: 266,610		
Non-trainable params: 0		



```
# compile model
```

```
model.compile(loss="sparse_categorical_crossentropy",  
              optimizer="sgd",  
              metrics=["accuracy"])
```

```
# train the model
```


```
history = model.fit(X_train, y_train, epochs=10, validation_data=(X_valid, y_valid))
```

```
# evaluate the model
```

```
model.evaluate(X_test, y_test)
```

```
# get some predictions
```

```
y_proba = model.predict(X_new)
```



Callbacks

- Callbacks allow us to modify the default behavior of the training loop.

```
# save best model during training
```

```
checkpoint_cb = keras.callbacks.ModelCheckpoint("my_keras_model.h5",  
save_best_only=True)
```

```
history = model.fit(X_train, y_train, epochs=10,  
                    validation_data=(X_valid, y_valid),  
                    callbacks=[checkpoint_cb])
```

```
# rollback to best model
```

```
model = keras.models.load_model("my_keras_model.h5")
```

```
# custom callbacks
```

```
class PrintValTrainRatioCallback(keras.callbacks.Callback):  
    def on_epoch_end(self, epoch, logs):  
        print("\nval/train: {:.2f}".format(logs["val_loss"] / logs["loss"]))
```

```
val_train_ratio_cb = PrintValTrainRatioCallback()
```

```
history = model.fit(X_train, y_train, epochs=3,  
                    validation_data=(X_valid, y_valid),  
                    callbacks=[val_train_ratio_cb])
```

```
Epoch 1/3  
280/363 [=====>.....] - ETA: 0s - loss: 0.3830  
val/train: 0.94  
363/363 [=====] - 0s 752us/step - loss: 0.3913 - val_loss: 0.3677  
Epoch 2/3  
356/363 [=====>.] - ETA: 0s - loss: 0.3927  
val/train: 0.94  
363/363 [=====] - 0s 782us/step - loss: 0.3903 - val_loss: 0.3676  
Epoch 3/3  
274/363 [=====>.....] - ETA: 0s - loss: 0.3905  
val/train: 0.96  
363/363 [=====] - 0s 749us/step - loss: 0.3890 - val_loss: 0.3744
```

Keras

- Great high level API with abstraction
- Attractive for less experienced practitioners
- Useful for people outside the CS field
- Tensorflow ecosystem (hardware, deployment, ...)
- Some flexibility if necessary, but for non-standards problems we may need to go lower level -> **Tensorflow !**

Tensorflow

```
from tensorflow.keras.layers import Dense, Flatten, Conv2D
from tensorflow.keras import Model

# Build the dataset

train_ds = tf.data.Dataset.from_tensor_slices(
    (x_train, y_train)).shuffle(10000).batch(32)

test_ds = tf.data.Dataset.from_tensor_slices((x_test, y_test)).batch(32)

# optimizer, loss and metrics

loss_object = tf.keras.losses.SparseCategoricalCrossentropy()

optimizer = tf.keras.optimizers.Adam()

train_loss = tf.keras.metrics.Mean(name='train_loss')
train_accuracy = tf.keras.metrics.SparseCategoricalAccuracy(name='train_accuracy')

test_loss = tf.keras.metrics.Mean(name='test_loss')
test_accuracy = tf.keras.metrics.SparseCategoricalAccuracy(name='test_accuracy')
```

```
# Define a network
```

```
class MyModel(Model):  
    def __init__(self):  
        super(MyModel, self).__init__()  
        self.conv1 = Conv2D(32, 3, activation='relu')  
        self.flatten = Flatten()  
        self.d1 = Dense(128, activation='relu')  
        self.d2 = Dense(10, activation='softmax')  
  
    def call(self, x):  
        x = self.conv1(x)  
        x = self.flatten(x)  
        x = self.d1(x)  
        return self.d2(x)  
  
model = MyModel()
```



```

@tf.function
def train_step(images, labels):
    with tf.GradientTape() as tape:
        predictions = model(images)
        loss = loss_object(labels, predictions)
    gradients = tape.gradient(loss, model.trainable_variables)
    optimizer.apply_gradients(zip(gradients, model.trainable_variables))

    train_loss(loss)
    train_accuracy(labels, predictions)

@tf.function
def test_step(images, labels):
    predictions = model(images)
    t_loss = loss_object(labels, predictions)

    test_loss(t_loss)
    test_accuracy(labels, predictions)

EPOCHS = 5

for epoch in range(EPOCHS):
    for images, labels in train_ds:
        train_step(images, labels)

    for test_images, test_labels in test_ds:
        test_step(test_images, test_labels)

    template = 'Epoch {}, loss: {}, acc: {}, test_loss: {}, test_acc: {}'
    print(template.format(epoch+1,
                           train_loss.result(),
                           train_accuracy.result()*100,
                           test_loss.result(),
                           test_accuracy.result()*100))

    # reset metrics for next epoch.
    train_loss.reset_states()
    train_accuracy.reset_states()
    test_loss.reset_states()
    test_accuracy.reset_states()

```

Tensorflow

- Low level API (ideal for research)
- Need to define your own dataset
- Need to define your own training loop
- More control and flexibility
- Great ecosystem and production ready tools.
- Still some details that make it confusing sometimes...





- Most used framework in research and competitions
- Developed by Facebook
- No official High-level API (like Keras), but a good option is Fast.ai (plus really good education content).
- Low level API
- Catching up on the production side (torchscript, torch serve, ...)

Pytorch

```
# import pytorch

import torch
import torchvision

# download data

import torchvision.transforms as transforms

train_dataset = torchvision.datasets.MNIST('./data', train=True, download=True,
                                           transform=transforms.Compose([
                                               transforms.ToTensor(),
                                               transforms.Normalize((0.1307,), (0.3081,))
                                           ]))

test_dataset = torchvision.datasets.MNIST('./data', train=False,
                                           transform=transforms.Compose([
                                               transforms.ToTensor(),
                                               transforms.Normalize((0.1307,), (0.3081,))
                                           ]))

train_loader = torch.utils.data.DataLoader(train_dataset, batch_size=32,
                                           shuffle=True)
test_loader = torch.utils.data.DataLoader(test_dataset, batch_size=10000)
```


2

2

5

5

7

7

5

5

9

4

5

5

4

4

9

9

4

4

9

9

3

3

8

8

2

2

0

0

0

0

2

2

7

7

0

0

1

1

1

1

2

2

9

9

5

5

9

9

6

6

3

3

9

9

9

9

```
# define the network
```

```
model = torch.nn.Sequential(  
    torch.nn.Linear(28*28, 100),  
    torch.nn.ReLU(),  
    torch.nn.Linear(100, 10)  
)
```

```
# optimizer and loss function
```

```
optimizer = torch.optim.SGD(model.parameters(), lr=0.1)  
loss = torch.nn.CrossEntropyLoss()
```

```

# train

def fit(model, train_loader, test_loader, optimizer, loss, epochs = 5):
    for epoch in range(1, epochs+1):
        # train
        model.train() # VERY IMPORTANT !
        train_loss, train_acc = [], 0
        for imgs, labels in tqdm(train_loader):
            optimizer.zero_grad() # VERY IMPORTANT !
            output = model(imgs.reshape(imgs.shape[0], 28*28))
            l = loss(output, labels)
            l.backward()
            optimizer.step()
            train_acc += (torch.argmax(output, axis=1) == labels).sum()
            train_loss.append(l.item())
        # eval
        model.eval() # VERY IMPORTANT !
        test_loss, test_acc = [], 0
        with torch.no_grad():
            for imgs, labels in tqdm(test_loader):
                output = model(imgs.reshape(imgs.shape[0], 28*28))
                l = loss(output, labels)
                test_acc += (torch.argmax(output, axis=1) == labels).sum()
                test_loss.append(l.item())
        print(f'Epoch: {epoch}/{epochs}\n Train loss:
{np.mean(train_loss):.4f} Test loss: {np.mean(test_loss):.4f} Train Acc:
{train_acc}/{len(train_dataset)} ({100*train_acc/len(train_dataset)} %)
Test Acc: {test_acc}/{len(test_dataset)} ({100*test_acc/
len(test_dataset)} %)')

```

```
# custom datasets

from torch.utils.data import Dataset

class MNISTDataset(Dataset):
    def __init__(self, images, labels):
        self.images = images
        self.labels = labels

    def __len__(self):
        return len(self.labels)

    def __getitem__(self, ix):
        img = torch.from_numpy(self.images[ix] / 255).float()
        return img, torch.tensor(self.labels[ix]).long()

train_dataset = MNISTDataset(X_train, y_train)
test_dataset = MNISTDataset(X_test, y_test)

# get first sample

img, label = train_dataset[0]
```

```
# custom networks

import torch.nn.functional as F

class Net(torch.nn.Module):

    def __init__(self):
        super(Net, self).__init__()
        self.fc1 = torch.nn.Linear(784, 100)
        self.fc2 = torch.nn.Linear(100, 10)

    def forward(self, x):
        x = self.fc1(x)
        x = F.relu(x)
        x = self.fc2(x)
        return x

net = Net()
```



```
# train on GPU
```

```
device = "cuda" if torch.cuda.is_available() else "cpu"
```

```
def fit_gpu(model, train_loader, test_loader, optimizer, loss, device, epochs=10):
```

```
    model.to(device) # COPY MODEL TO GPU
```

```
    for epoch in range(1, epochs+1):
```

```
        # train
```

```
        model.train() # VERY IMPORTANT !
```

```
        train_loss, train_acc = [], 0
```

```
        for imgs, labels in tqdm(train_loader):
```

```
            imgs, labels = imgs.to(device), labels.to(device) # COPY DATA TO GPU
```

```
            optimizer.zero_grad() # VERY IMPORTANT !
```

```
            output = model(imgs.reshape(imgs.shape[0], 28*28))
```

```
            l = loss(output, labels)
```

```
            l.backward()
```

```
            optimizer.step()
```

```
            train_acc += (torch.argmax(output, axis=1) == labels).sum()
```

```
            train_loss.append(l.item())
```

```
        # eval
```

```
        model.eval() # VERY IMPORTANT !
```

```
        test_loss, test_acc = [], 0
```

```
        with torch.no_grad():
```

```
            for imgs, labels in tqdm(test_loader):
```

```
                imgs, labels = imgs.to(device), labels.to(device) # COPY DATA TO GPU
```

```
                output = model(imgs.reshape(imgs.shape[0], 28*28))
```

```
                l = loss(output, labels)
```

```
                test_acc += (torch.argmax(output, axis=1) == labels).sum()
```

```
                test_loss.append(l.item())
```

```
        print(f'Epoch: {epoch}/{epochs}\n Train loss: {np.mean(train_loss):.4f}
```

```
Test loss: {np.mean(test_loss):.4f} Train Acc: {train_acc}/{len(train_dataset)}
```

```
({100*train_acc/len(train_dataset)} %) Test Acc: {test_acc}/{len(test_dataset)}
```

```
({100*test_acc/len(test_dataset)} %)' )
```

```
# save model
```

```
torch.save(net.state_dict(), "my_model_params.pth")
```

```
torch.save(net, "my_model.pth")
```

```
scripted_model = torch.jit.script(net)
```

```
scripted_model.save("my_scripted_model.pth")
```

```
# load model
```

```
net.load_state_dict(torch.load("my_model_params.pth"))
```

```
net = torch.load("my_model.pth")
```

```
net = torch.jit.load("my_scripted_model.pth")
```

Pytorch

- Low level API (ideal for research)
- Need to define your own datasets
- Need to define your own training loop
- More control and flexibility
- Great ecosystem.
- Still behind on production ready tools.

Other Frameworks & Interoperability


Other Frameworks

- MXNet (Amazon)
- CNTK (Microsoft)
- CoreML (Apple - iOS)
- Chainer
- ...

Interoperability

- Standardize neural network representations
- Train in one framework, optimize in another, deploy in another.
- Opportunity for optimization (pruning, quantization, ...)
- Language agnostic
- Example: ONNX (<https://onnx.ai/>)

<https://colab.research.google.com/github/sensioai/dl/blob/master/frameworks/frameworks.ipynb>



🔑 WOULD YOU LIKE TO KNOW MORE?

<https://keras.io/>

<https://www.tensorflow.org/>

<https://www.fast.ai/>

<https://pytorch.org/>

Practice with this exercise



<https://colab.research.google.com/github/sensioai/dl/blob/master/frameworks/exercise.ipynb>

The background of the slide features a bokeh effect with numerous out-of-focus circles in shades of yellow, orange, and blue against a dark navy blue background. The circles vary in size and brightness, creating a soft, glowing texture.

DEEP LEARNING FRAMEWORKS