


The background of the slide features a bokeh effect with numerous out-of-focus circles in shades of yellow, orange, and blue against a dark navy blue background. The circles vary in size and opacity, creating a soft, glowing effect.

The Multilayer Perceptron


Contents

- Introduction
- Regression
- Classification
- Our own MLP framework



Introduction


Perceptron



The Perceptron computes a weighted sum of its inputs and then applies an activation function.

$$\hat{y} = f(w_0 + w_1x_1 + \dots + w_mx_m)$$

Multilayer Perceptron



The Multilayer Perceptron (MLP) stacks Perceptrons in sequential layers, feeding the outputs of one layer to the inputs of the next layer. In the case of a 2 layer MLP,

$$\hat{y} = f_2 \left(w_0^2 + w_1^2 h_1 + \dots + w_{m_h}^2 h_{m_h} \right)$$

$$h = f_1 \left(w_0^1 + w_1^1 x_1 + \dots + w_m^1 x_m \right)$$


Training an MLP

Gradient descent on the perceptron

- Compute the output, \hat{y}
- Compute the gradient of the loss function wrt the parameters, $\frac{\partial l}{\partial w}$
- Update parameters, $w \leftarrow w - \eta \frac{\partial l}{\partial w}$
- Repeat until convergence

Training an MLP

Backpropagation



Regression


```

class MLP():
    def __init__(self, D_in, H, D_out):
        self.w1, self.b1 = np.random.normal(loc=0.0,
                                              scale=np.sqrt(2/(D_in+H)),
                                              size=(D_in, H)), np.zeros(H)
        self.w2, self.b2 = np.random.normal(loc=0.0,
                                              scale=np.sqrt(2/(H+D_out)),
                                              size=(H, D_out)), np.zeros(D_out)

        self.loss = mse
        self.grad_loss = grad_mse

    def __call__(self, x):
        self.h = np.dot(x, self.w1) + self.b1
        y_hat = np.dot(self.h, self.w2) + self.b2
        return self.final_activation(y_hat)


    def final_activation(self, x):
        return x

def mse(output, target):
    return 0.5*(output - target)**2


def grad_mse(output, target):
    return (output - target)

def fit(self, X, Y, epochs = 100, lr = 0.001):
    for e in range(epochs):
        for x, y in zip(X, Y):
            # add batch dimension
            x = x[None,:]
            y_pred = self(x)
            # loss function
            loss = self.loss(y_pred, y).mean()
            # Backprop
            # dl/dy
            dldy = self.grad_loss(y_pred, y)
            # dl/dw2 = dl/dy * dy/dw2
            grad_w2 = np.dot(self.h.T, dldy)
            grad_b2 = dldy.mean(axis=0)*self.h.shape[0]
            # dl/dh = dl/dy * dy/dh
            dldh = np.dot(dldy, self.w2.T)*self.h
            # dl/dw1 = dl/dy * dy/dh * dh/dw1
            grad_w1 = np.dot(x.T, dldh)
            grad_b1 = dldh.mean(axis=0)*x.shape[0]
            # Update (GD)
            self.w1 = self.w1 - lr * grad_w1
            self.b1 = self.b1 - lr * grad_b1
            self.w2 = self.w2 - lr * grad_w2
            self.b2 = self.b2 - lr * grad_b2

```



Polynomial Regression



A linear combination of linear functions is a linear function. We need non linearity !

```

def relu(x):
    return np.maximum(0, x)


def reluPrime(x):
    return x > 0

class MLPrelu(MLP):
    def __call__(self, x):
        self.h = relu(np.dot(x, self.w1) + self.b1)
        y_hat = np.dot(self.h, self.w2) + self.b2
        return self.final_activation(y_hat)

    def fit(self, X, Y, epochs = 100, lr = 0.001):
        for e in range(epochs):
            for x, y in zip(X, Y):
                x = x[None,:]
                y_pred = self(x)
                loss = self.loss(y_pred, y).mean()
                # Backprop
                dldy = self.grad_loss(y_pred, y)
                grad_w2 = np.dot(self.h.T, dldy)
                grad_b2 = dldy
                dldh = np.dot(dldy, self.w2.T)*reluPrime(self.h)
                grad_w1 = np.dot(x.T, dldh)
                grad_b1 = dldh
                # Update (GD)
                self.w1 = self.w1 - lr * grad_w1
                self.b1 = self.b1 - lr * grad_b1
                self.w2 = self.w2 - lr * grad_w2
                self.b2 = self.b2 - lr * grad_b2

```


Epoch 50



Binary Classification