


The background of the slide features a bokeh effect with numerous out-of-focus circles in shades of yellow, orange, and blue against a dark navy blue background. The circles vary in size and opacity, creating a soft, glowing texture.

The Multilayer Perceptron


Contents

- Introduction
- Regression
- Classification
- Our own MLP framework



Introduction


Perceptron



The Perceptron computes a weighted sum of its inputs and then applies an activation function.

$$\hat{y} = f(w_0 + w_1x_1 + \dots + w_mx_m)$$

Multilayer Perceptron



The Multilayer Perceptron (MLP) stacks Perceptrons in sequential layers, feeding the outputs of one layer to the inputs of the next layer. In the case of a 2 layer MLP,

$$\hat{y} = f_2 \left(w_0^2 + w_1^2 h_1 + \dots + w_{m_h}^2 h_{m_h} \right)$$

$$h = f_1 \left(w_0^1 + w_1^1 x_1 + \dots + w_m^1 x_m \right)$$


Training an MLP

Gradient descent on the perceptron


- Compute the output, \hat{y}
- Compute the gradient of the loss function w.r.t the parameters, $\partial l / \partial w$
- Update parameters, $w \leftarrow w - \eta \frac{\partial l}{\partial w}$
- Repeat until convergence

Training an MLP

Backpropagation



Backpropagation



Assuming MSE loss function and linear activation function at the output

$$\frac{\partial l}{\partial w^2} = \frac{\partial l}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial w^2} = (\hat{y} - y)h$$

$$\frac{\partial l}{\partial w^1} = \frac{\partial l}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial h} \frac{\partial h}{\partial w^1} = (\hat{y} - y)w^2 f'_1(w^1 x)x$$

Regression

```

class MLP():
    def __init__(self, D_in, H, D_out):
        self.w1, self.b1 = np.random.normal(loc=0.0,
                                              scale=np.sqrt(2/(D_in+H)),
                                              size=(D_in, H)), np.zeros(H)
        self.w2, self.b2 = np.random.normal(loc=0.0,
                                              scale=np.sqrt(2/(H+D_out)),
                                              size=(H, D_out)), np.zeros(D_out)

        self.loss = mse
        self.grad_loss = grad_mse

    def __call__(self, x):
        self.h = np.dot(x, self.w1) + self.b1
        y_hat = np.dot(self.h, self.w2) + self.b2
        return self.final_activation(y_hat)


    def final_activation(self, x):
        return x

def mse(output, target):
    return 0.5*(output - target)**2


def grad_mse(output, target):
    return (output - target)

def fit(self, X, Y, epochs = 100, lr = 0.001):
    for e in range(epochs):
        for x, y in zip(X, Y):
            # add batch dimension
            x = x[None,:]
            y_pred = self(x)
            # loss function
            loss = self.loss(y_pred, y).mean()
            # Backprop
            # dl/dy
            dldy = self.grad_loss(y_pred, y)
            # dl/dw2 = dl/dy * dy/dw2
            grad_w2 = np.dot(self.h.T, dldy)
            grad_b2 = dldy
            # dl/dh = dl/dy * dy/dh
            dldh = np.dot(dldy, self.w2.T)
            # dl/dw1 = dl/dy * dy/dh * dh/dw1
            grad_w1 = np.dot(x.T, dldh)
            grad_b1 = dldh
            # Update (GD)
            self.w1 = self.w1 - lr * grad_w1
            self.b1 = self.b1 - lr * grad_b1
            self.w2 = self.w2 - lr * grad_w2
            self.b2 = self.b2 - lr * grad_b2

```



Polynomial Regression



A linear combination of linear functions is a linear function. We need non linearity !

```

def relu(x):
    return np.maximum(0, x)


def reluPrime(x):
    return x > 0

class MLPrelu(MLP):
    def __call__(self, x):
        self.h = relu(np.dot(x, self.w1) + self.b1)
        y_hat = np.dot(self.h, self.w2) + self.b2
        return self.final_activation(y_hat)

    def fit(self, X, Y, epochs = 100, lr = 0.001):
        for e in range(epochs):
            for x, y in zip(X, Y):
                x = x[None,:]
                y_pred = self(x)
                loss = self.loss(y_pred, y).mean()
                # Backprop
                dldy = self.grad_loss(y_pred, y)
                grad_w2 = np.dot(self.h.T, dldy)
                grad_b2 = dldy
                dldh = np.dot(dldy, self.w2.T)*reluPrime(self.h)
                grad_w1 = np.dot(x.T, dldh)
                grad_b1 = dldh
                # Update (GD)
                self.w1 = self.w1 - lr * grad_w1
                self.b1 = self.b1 - lr * grad_b1
                self.w2 = self.w2 - lr * grad_w2
                self.b2 = self.b2 - lr * grad_b2

```


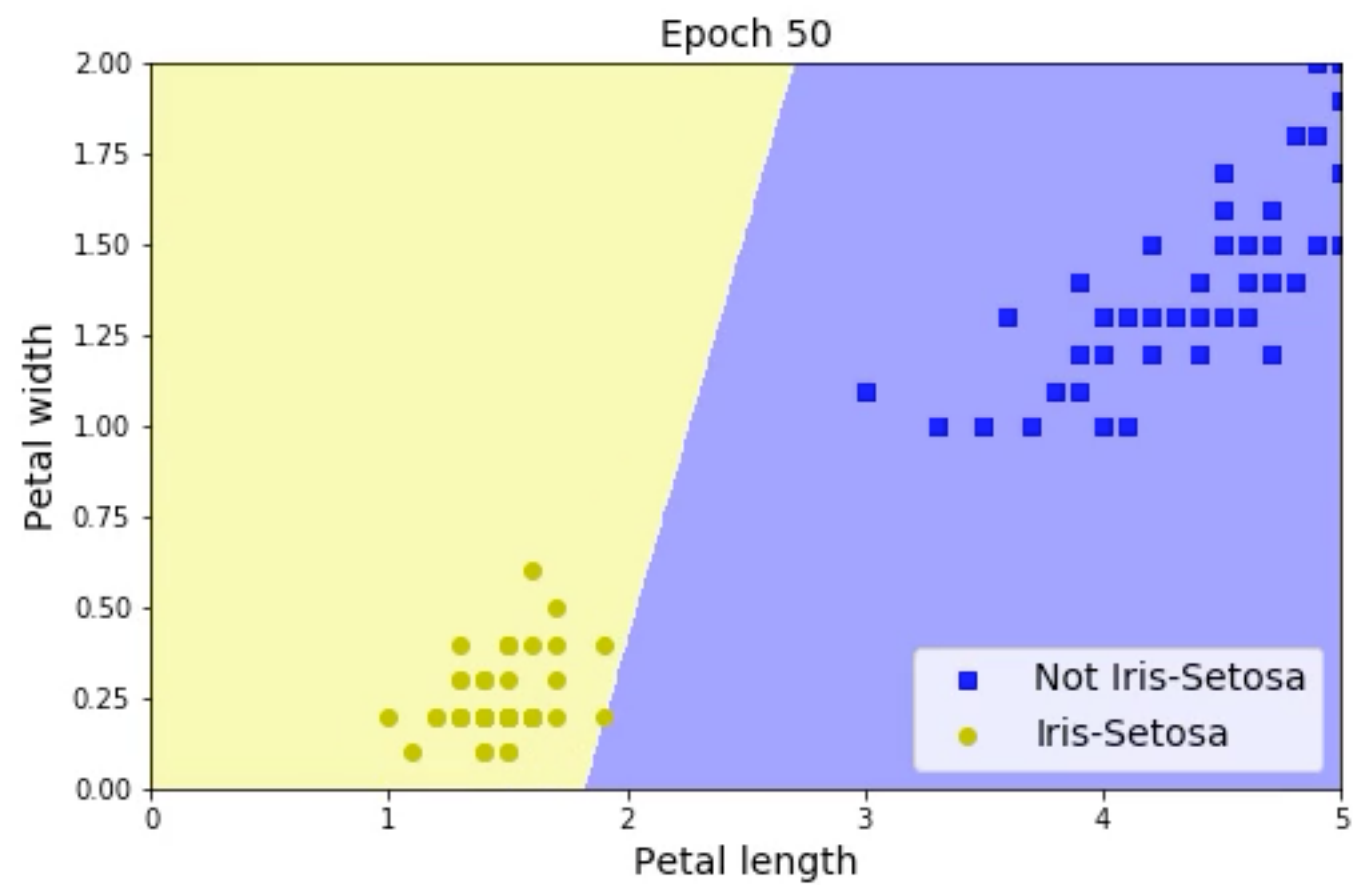

Epoch 50



Binary Classification

```
class MLPreluBinaryClass(MLPrelu):  
    def final_activation(self, x):  
        return x > 0
```

```
class MLPrelu(MLP):  
    def __call__(self, x):  
        self.h = relu(np.dot(x, self.w1) + self.b1)  
        y_hat = np.dot(self.h, self.w2) + self.b2  
        return self.final_activation(y_hat)  
  
    def fit(self, X, Y, epochs = 100, lr = 0.001):  
        for e in range(epochs):  
            for x, y in zip(X, Y):  
                x = x[None, :]  
                y_pred = self(x)  
                loss = self.loss(y_pred, y).mean()  
                # Backprop  
                dldy = self.grad_loss(y_pred, y)  
                grad_w2 = np.dot(self.h.T, dldy)  
                grad_b2 = dldy  
                dldh = np.dot(dldy, self.w2.T) * reluPrime(self.h)  
                grad_w1 = np.dot(x.T, dldh)  
                grad_b1 = dldh  
                # Update (GD)  
                self.w1 = self.w1 - lr * grad_w1  
                self.b1 = self.b1 - lr * grad_b1  
                self.w2 = self.w2 - lr * grad_w2  
                self.b2 = self.b2 - lr * grad_b2
```



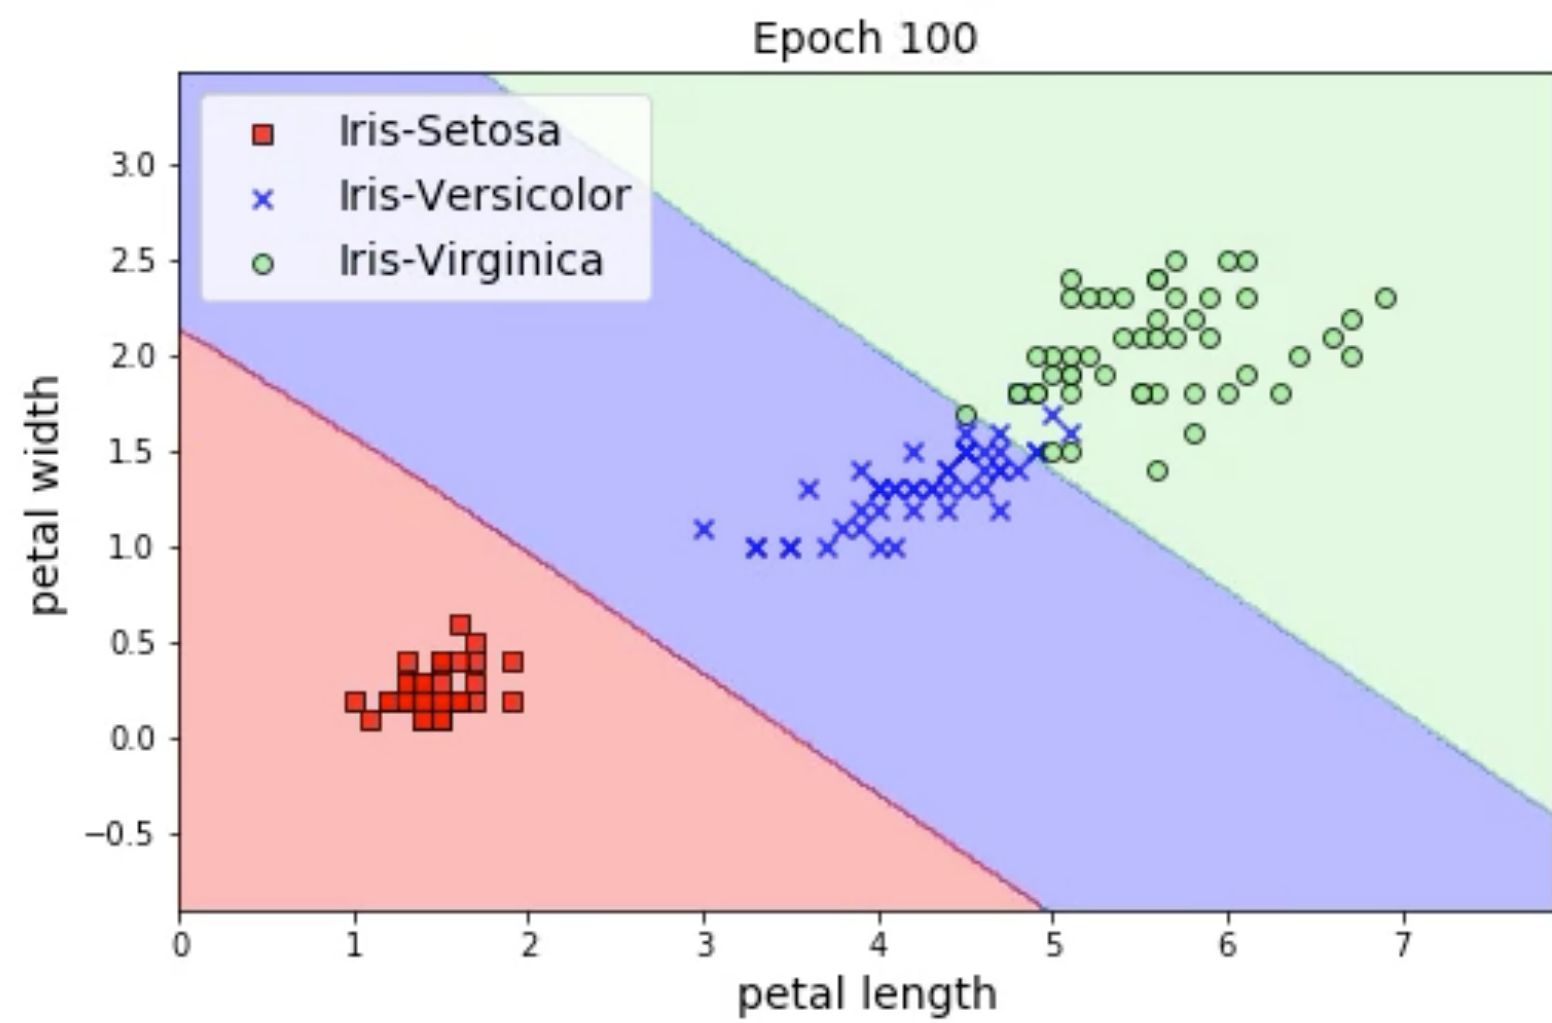
Multiclass Classification


```
def crossentropy(output, target):
    logits = output[np.arange(len(output)), target]
    entropy = - logits + np.log(np.sum(np.exp(output), axis=-1))
    return entropy

def softmax(x):
    return np.exp(x) / np.exp(x).sum(axis=-1, keepdims=True)

def grad_crossentropy(output, target):
    answers = np.zeros_like(output)
    answers[np.arange(len(output)), target] = 1
    return (- answers + softmax(output)) / output.shape[0]

class MLPreluClass(MLPrelu):
    def __init__(self, D_in, H, D_out):
        super().__init__(D_in, H, D_out)
        self.loss = crossentropy
        self.grad_loss = grad_crossentropy
```



What if now ... ?

- We want to use a different activation function (not relu) ?
- We want to use more than 2 layers (3, 4, 5...) ?
- We want to use a different optimization algorithm (not SGD) ?
- We need more flexibility... we need an MLP framework !

Building an MLP framework

Pytorch-like API

```
D_in, H, D_out = 2, 3, 2

mlp = MLP([
    Linear(D_in, H),
    ReLU(),
    Linear(H, D_out)
])

optimizer = SGD(mlp, lr=0.1)
loss = CrossEntropy(mlp)

epochs = 100
for e in range(epochs):
    for x, y in zip(X, Y):
        y_pred = mlp(x)
        loss(y_pred, y)
        loss.backward()
        optimizer.update()
```


MLP

```
class MLP:
    def __init__(self, layers):
        self.layers = layers

    def __call__(self, x):
        for layer in self.layers:
            x = layer(x)
        return x
```

Layers

```
class Layer():  
    def __init__(self):  
        self.params = []  
        self.grads = []  
  
    def __call__(self, x):  
        return x  
  
    def backward(self, grad):  
        return grad  
  
    def update(self, params):  
        return
```

```
class Linear(Layer):
    def __init__(self, d_in, d_out):
        self.w = np.random.normal(loc=0.0,
                                    scale=np.sqrt(2/(d_in+d_out)),
                                    size=(d_in, d_out))

        self.b = np.zeros(d_out)

    def __call__(self, x):
        self.x = x
        self.params = [self.w, self.b]
        return np.dot(x, self.w) + self.b

    def backward(self, grad_output):
        grad = np.dot(grad_output, self.w.T)
        self.grad_w = np.dot(self.x.T, grad_output)
        self.grad_b = grad_output.mean(axis=0)*self.x.shape[0]
        self.grads = [self.grad_w, self.grad_b]
        return grad

    def update(self, params):
        self.w = params[0]
        self.b = params[1]
```

```
class ReLU(Layer):  
    def __call__(self, x):  
        self.x = x  
        return np.maximum(0, x)  
  
    def backward(self, grad_output):  
        grad = self.x > 0  
        return grad_output*grad
```

Optimizers

```
class SGD():
    def __init__(self, net, lr):
        self.net = net
        self.lr = lr

    def update(self):
        for layer in self.net.layers:
            layer.update([
                params - self.lr*grads
                for params, grads in zip(layer.params, layer.grads)
            ])
```



Losses

```
class CrossEntropy():
    def __init__(self, net):
        self.net = net


    def __call__(self, output, target):
        self.output, self.target = output, target
        logits = output[np.arange(len(output)), target]
        loss = - logits + np.log(np.sum(np.exp(output), axis=-1))
        loss = loss.mean()
        return loss

    def grad_crossentropy(self):
        answers = np.zeros_like(self.output)
        answers[np.arange(len(self.output)), self.target] = 1
        return (- answers + softmax(self.output)) / self.output.shape[0]

    def backward(self):
        grad = self.grad_crossentropy()
        for layer in reversed(self.net.layers):
            grad = layer.backward(grad)
```



<https://colab.research.google.com/github/sensioai/nbs/blob/master/mlp/mlp.ipynb>



🔑 WOULD YOU LIKE TO KNOW MORE?

<https://playground.tensorflow.org/>

Let's code !



<https://colab.research.google.com/github/sensioai/nbs/blob/master/mlp/exercise.ipynb>

The background of the slide features a bokeh effect with numerous out-of-focus circles in shades of yellow, orange, and blue against a dark grey or black background. The circles vary in size and brightness, creating a soft, glowing texture.

The Multilayer Perceptron