


The background of the image is a dark, textured surface covered with numerous out-of-focus light circles, known as bokeh. These circles are primarily in shades of warm orange and yellow, with some cooler blue and teal tones interspersed, particularly towards the right side. The circles vary in size and brightness, creating a dynamic and visually appealing pattern.

NATURAL LANGUAGE PROCESSING

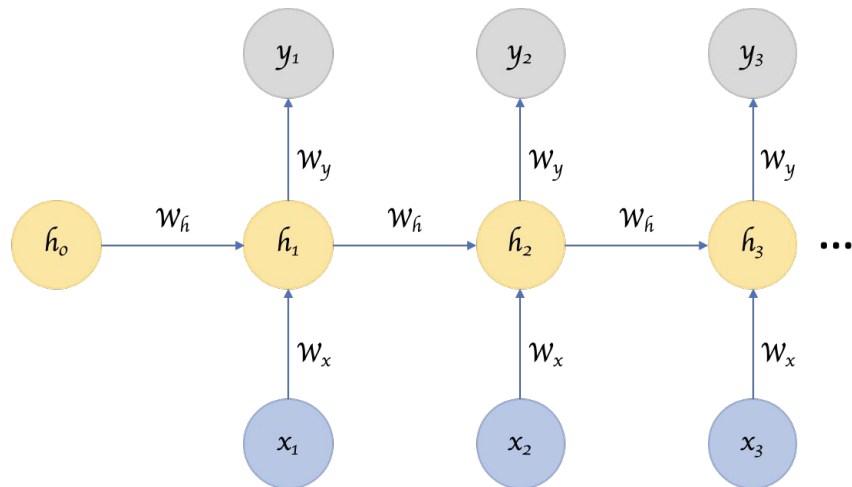
CONTENTS

- RNN review
- Natural Language Processing
- Character RNN
 - Tokenization
 - Embeddings
 - Generating text one character at a time
- Sentiment Analysis
 - Torchtext
 - Bi-directional RNNs
 - Transfer learning
 - Text Classification
- Seq2Seq models
 - Encoder-Decoder architecture
 - Machine translation
- Attention Mechanisms
 - Transformers
 - BERT for text classification



RNN Review

RNN



$$y_t = \mathbf{W}_y \mathbf{h}_t = \mathbf{W}_y f(\mathbf{W}_x \mathbf{x}_t + \mathbf{W}_h \mathbf{h}_{t-1})$$

```
class RNN(torch.nn.Module):
    def __init__(self, n_in=50, n_out=1):
        super().__init__()
        self.rnn = torch.nn.RNN(input_size=1, hidden_size=20, num_layers=2, batch_first=True)
        self.fc = torch.nn.Linear(20, 1)

    def forward(self, x):
        x, h = self.rnn(x)
        # get the last output
        x = self.fc(x[:, -1])
        return x
```



Use LSTM or GRU for better results.

The background of the slide is a dark blue to purple gradient, overlaid with numerous out-of-focus, glowing circles in shades of light blue and cyan, creating a bokeh effect.

Natural Language Processing

Natural Language Processing

- Language comprehension (virtual assistants such as Siri, Alexa, ...)
- Machine translation (Google Translate, ...)
- Text generation (language modeling, summarization, question answering, ...)
- Text classification (sentiment analysis, identify hate speech on social media, ...)
- Text-to-speech (generate audio from text) and Speech-to-text
- Image captioning and OCR

⚡ NLP is a very active field at the moment, new huge architectures (Transformers) and training techniques (language modeling on big unsupervised datasets) are providing excellent results improving SOTA by large margin on almost every task. See for example: <https://openai.com/blog/openai-api/>

Character RNN

CharRNN

- Generate text, one character at a time (<https://github.com/karpathy/char-rnn>)

AUTOLYCUS:

This is a merry ballad, but a very pretty one.

MOPSA:

Let's have some merry ones.

AUTOLYCUS:

Why, this is a passing merry one and goes to the tune of 'Two maids wooing a man;' there's scarce a maid westward but she sings it; 'tis in request, I can tell you.

MOPSA:

We can both sing it: if thou'lt bear a part, thou shalt hear; 'tis in three parts.

DORCAS:

We had the tune on't a month ago.

AUTOLYCUS:

I can bear my part; you must know 'tis my occupation; have at it with you.

Tokenization


We need to transform each character to a number.

```
import string

class Tokenizer():
    def __init__(self):
        self.all_characters = string.printable
        self.n_characters = len(self.all_characters)

    def text_to_seq(self, string):
        seq = []
        for c in range(len(string)):
            seq.append(self.all_characters.index(string[c]))
        return seq

    def seq_to_text(self, seq):
        text = ''
        for c in range(len(seq)):
            text += self.all_characters[seq[c]]
        return text
```



0123456789abcdefghijklmnopqrstuvwxyz
zABCDEFGHIJKLMNOPQRSTUVWXYZ!"#\$%&'(
)*+,-./:;<=>?@[\\]^_`{|}~
\t\n\r\x0b\x0c'

```
tokenizer.text_to_seq('abcDEF')
> [10, 11, 12, 39, 40, 41]

tokenizer.seq_to_text([10, 11, 12])
> 'abc'
```

Creating text windows

```
def windows(text, window_size = 100):  
    start_index = 0  
    end_index = len(text) - window_size  
    text_windows = []  
    while start_index < end_index:  
        text_windows.append(text[start_index:start_index+window_size+1])  
        start_index += 1  
    return text_windows
```

```
> ['First Citizen:\nBefore we proceed any further, hear me speak.\n\nAll:\nSpeak, speak.\n\nFirst  
Citizen:\nYou ',
```

```
'irst Citizen:\nBefore we proceed any further, hear me speak.\n\nAll:\nSpeak, speak.\n\nFirst  
Citizen:\nYou a',
```


```
'rst Citizen:\nBefore we proceed any further, hear me speak.\n\nAll:\nSpeak, speak.\n\nFirst  
Citizen:\nYou ar']
```

Dataset

```
class CharRNNDataset(torch.utils.data.Dataset):  
    def __init__(self, text_encoded_windows, train=True):  
        self.text = text_encoded_windows  
        self.train = train  
  
    def __len__(self):  
        return len(self.text)  
  
    def __getitem__(self, ix):  
        if self.train:  
            return torch.tensor(self.text[ix][:-1]), torch.tensor(self.text[ix][-1])  
        return torch.tensor(self.text[ix])
```

Embeddings


- The tokenizer converts each character into a number (0-99).
- We need to transform each number (class) to either:
 - Embeddings
 - One-hot encoding




	1	2	3	4	5	6	7	8	9
man	1	0	0	0	0	0	0	0	0
woman	0	1	0	0	0	0	0	0	0
boy	0	0	1	0	0	0	0	0	0
girl	0	0	0	1	0	0	0	0	0
prince	0	0	0	0	1	0	0	0	0
princess	0	0	0	0	0	1	0	0	0
queen	0	0	0	0	0	0	1	0	0
king	0	0	0	0	0	0	0	1	0
monarch	0	0	0	0	0	0	0	0	1

	Femininity	Youth	Royalty
Man	0	0	0
Woman	1	0	0
Boy	0	1	0
Girl	1	1	0
Prince	0	1	1
Princess	1	1	1
Queen	1	0	1
King	0	0	1
Monarch	0.5	0.5	1


Vocabulary:
Man, woman, boy,
girl, prince,
princess, queen,
king, monarch



Male-Female



Verb tense



Country-Capital

Generating text

```
X_new = "With eyes wide open; standing, speaking, movin"
```

```
for i in range(1000):  
    X_new_encoded = tokenizer.text_to_seq(X_new[-100:])  
    y_pred = model.predict(X_new_encoded)  
    y_pred = torch.argmax(y_pred, axis=1)[0].item()  
    X_new += tokenizer.seq_to_text([y_pred])
```

```
> 'With eyes wide open; standing, speaking, movingment\nWhere should should should should should  
should should should should should should should should she shall she\nshall should should  
should should should should should should should should should should should she shall  
she\nshall should should should should should should should should should should should  
should should she shall she\nshall should should should should should should should should  
should should should should should she shall she\nshall should should should should should  
should should should should should should should should she shall she\nshall should should should  
should should should should should should should should should should should she shall she\nshall  
should should should should should should should should should should should should should  
she shall she\nshall should should should should should should should should should should should  
should should should she shall she\nshall should should should should should should sh'
```

```

X_new = "With eyes wide open; standing, speaking, movin"

temp= 0.8
for i in range(1000):
    X_new_encoded = tokenizer.text_to_seq(X_new[-100:])
    y_pred = model.predict(X_new_encoded)
    y_pred = y_pred.view(-1).div(temp).exp()
    top_i = torch.multinomial(y_pred, 1)[0]
    predicted_char = tokenizer.all_characters[top_i]
    X_new += predicted_char

```

With eyes wide open; standing, speaking, moving,
 Do the word, gain is the own spoke quick, and my heart
 Leates of his face.
 Come and all from me reday:
 Now not seem by sight?

Third Citizen:
 Now bring the life of aid the deprisonant contrees;
 For where to your repongine enough here stright
 Should sure them with lupken Cite best we weed should
 should not leave:
 What can here is thee he is my child and there is thee,
 Not son tears of me believe the best see purgal they she
 Twict in the noble with his words of stross
 Of that you cut make you to be confess,
 His stricks, but myself free so popurine of yourself.

LUCIO:
 Now, for her hand? that sing for the prince is speak
 Ot shall desire death, serve for that what we die me;
 Send you are passion vapeen and move princes me them,
 Which you that knows that from from your airful truester
 My lord, for when I see, speak, and i' the conspiracy,
 Luth as promised.

DUCHESS OF YORK:
 Now, with this heart, tyrant me shall the refore that you
 emstre:
 Gaunter, known her stept that is grace is the sheep

Sentiment Analysis

Sentiment Analysis

- Sentiment analysis is a case of text classification, where we want to assign a particular label to a piece of text (detect hate speech in social networks, user satisfaction on product reviews,)

Reviews

Average Rating: ★★★★★ (based on 3 reviews)

Showing 3 Reviews:

★★★★★	by Lucy on 5/10/2008	Great product, works better than expected I am so impressed that I'll be buying another. I will be recommending this to all my friends.
★★★★☆	by Customer on 5/9/2008	I wish I went with HDTV instead For me, the HDTV would have been better. The Plasma TV I bought was damaged right away, because I left it on overnight. The images burned into the screen. I should have been more careful.
★★★★★	by Joe on 9/10/2007	Great value - Great Quality Sanasonic introduced the SN-42PX60U with an MSRP of \$2,499 but it can easily be found around \$2,000 at one of several Internet retailers. The SN-42PX60U is not only an impressive plasma display, but it is the most affordable of the top-tier plasma televisions. Both Wamsung and Tushiba have 42" plasma televisions on the market with marginally lower prices, but they can hardly compete with the superior picture quality of the SN-42PX60U. Bioneer's PDP-4360HD has a very high-quality picture, but Bioneer also includes superfluous features such as a separate box which inflate the PDP-4360HD's price to roughly \$1,500 more than that of the SN-42PX60U. Perhaps the most interesting price comparison is between the SN-42PX60U and its predecessor, the SN-42PX50U. When it hit the market in 2005, the SN-42PX50U was heralded for its great value - it had an MSRP of \$2,999; \$500 more than the SN-42PX60U!

Rate this item



Torchtext

<https://pytorch.org/text/>

- Abstract and automate text data preprocessing (tokenization, data splitting, batching, ...)
- Includes some popular datasets, here we use IMDB movie reviews dataset.
- Customize to your needs.


```
import torch
import torchtext

TEXT = torchtext.data.Field(tokenize = 'spacy')
LABEL = torchtext.data.LabelField(dtype = torch.float)

train_data, test_data = torchtext.datasets.IMDB.splits(TEXT, LABEL)

MAX_VOCAB_SIZE = 10000
TEXT.build_vocab(train_data, max_size = MAX_VOCAB_SIZE)
LABEL.build_vocab(train_data)
```


Bidirectional RNNs



```
class BidirectionalRNN(RNN):
    def __init__(self, input_dim, embedding_dim=128, hidden_dim=128, output_dim=1, dropout=0.2,
pad_idx=0, bidirectional=True):
        super().__init__(input_dim, embedding_dim, hidden_dim, output_dim, dropout)
        self.rnn = torch.nn.GRU(input_size=embedding_dim,
                                hidden_size=hidden_dim,
                                num_layers=2,
                                dropout=dropout,
                                bidirectional=bidirectional)

    if bidirectional:
        self.fc = torch.nn.Linear(2*hidden_dim, output_dim)
```

Transfer Learning

- Use pre-trained embeddings

```
TEXT.build_vocab(train_data,
                  max_size = MAX_VOCAB_SIZE,
                  vectors = "glove.6B.100d",
                  unk_init = torch.Tensor.normal_)
LABEL.build_vocab(train_data)

pretrained_embeddings = TEXT.vocab.vectors

net.embedding.weight.data.copy_(pretrained_embeddings)

net.embedding.weight.data[TEXT.vocab.stoi[TEXT.unk_token]] = torch.zeros(100)
net.embedding.weight.data[TEXT.vocab.stoi[TEXT.pad_token]] = torch.zeros(100)
```

Text Classification

```
import spacy
nlp = spacy.load('en')

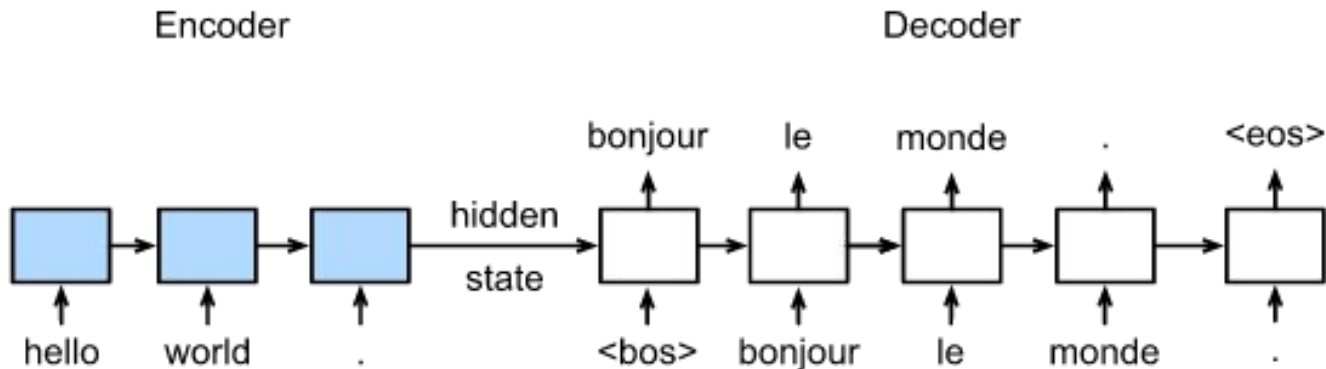
sentences = ["this film is terrible", "this film is great", "this film is good", "a waste of time"]
tokenized = [[tok.text for tok in nlp.tokenizer(sentence)] for sentence in sentences]
indexed = [[TEXT.vocab.stoi[_t] for _t in t] for t in tokenized]
tensor = torch.tensor(indexed).to(device).permute(1,0)
net.eval()
prediction = torch.sigmoid(net(tensor))

> [0.0732, 0.9613, 0.8762, 0.0119]
```

Sequence to sequence models

Sequence to sequence models

- Encoder-Decoder architecture.
- The encoder creates the initial state of the decoder.
- The decoder generates text one word at a time, using the output at each step as input for the next until termination.
- Used for machine translation, text summarization, ...
- If the encoder is a CNN, it can be used for image captioning, OCR, ...



EXERCISE !





https://pytorch.org/tutorials/intermediate/seq2seq_translation_tutorial.html

Attention Mechanisms

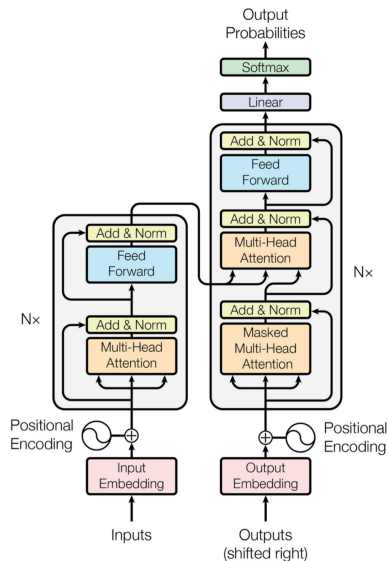
seq2seq with attention

- Attention mechanisms allows the network to focus on specific parts of its inputs.




Transformers


- A Transformer is a Neural Network architecture based on *Attention Mechanisms*
<https://arxiv.org/pdf/1706.03762.pdf>
- They are the SOTA today in NLP.



Scaled Dot-Product Attention



Multi-Head Attention



BERT for text classification

```
from transformers import BertTokenizer
```

pip install transformers

```
tokenizer = BertTokenizer.from_pretrained ('bert-base-uncased' )
```

```
max_input_length = tokenizer.max_model_input_sizes ['bert-base-uncased' ]
```

```
def tokenize_and_cut (sentence):
```

```
    tokens = tokenizer.tokenize (sentence)
```

```
    tokens = tokens [:max_input_length-2]
```

```
    return tokens
```

```
TEXT = torchtext.data.Field (batch_first = True,
```

```
                             use_vocab = False,
```

```
                             tokenize = tokenize_and_cut ,
```

```
                             preprocessing = tokenizer.convert_tokens_to_ids ,
```

```
                             init_token = tokenizer.cls_token_id ,
```

```
                             eos_token = tokenizer.sep_token_id ,
```

```
                             pad_token = tokenizer.pad_token_id ,
```

```
                             unk_token = tokenizer.unk_token_id )
```

```
LABEL = torchtext.data.LabelField (dtype = torch.float)
```




WOULD YOU LIKE TO KNOW MORE?

<https://github.com/sensioai/dl/tree/master/nlp>

Resources

Learn

- <https://www.youtube.com/watch?v=8rXD5-xhemo&list=PLoROMvodv4rOhcuXMZkNm7j3fVwBBY42z>
- <https://www.fast.ai/2019/07/08/fastai-nlp/>
- <https://www.youtube.com/watch?v=4jROIXH9Nvc>

Practice

- <https://pytorch.org/tutorials/index.html>
- <https://github.com/bentrevett/pytorch-sentiment-analysis>

The background of the image is a dark, textured surface covered with numerous out-of-focus light circles, known as bokeh. These circles are primarily in shades of warm orange and yellow, with some cooler blue and teal tones interspersed, particularly towards the top and right edges. The circles vary in size and brightness, creating a dynamic and visually appealing pattern.

NATURAL LANGUAGE PROCESSING