

The background of the image is a dark, textured surface covered with numerous out-of-focus light circles, known as bokeh. These circles vary in size and are primarily colored in warm tones of orange, yellow, and gold, with some cooler blue and teal circles interspersed. The overall effect is a soft, glowing, and abstract pattern.

# RECURRENT NEURAL NETWORKS

# CONTENTS

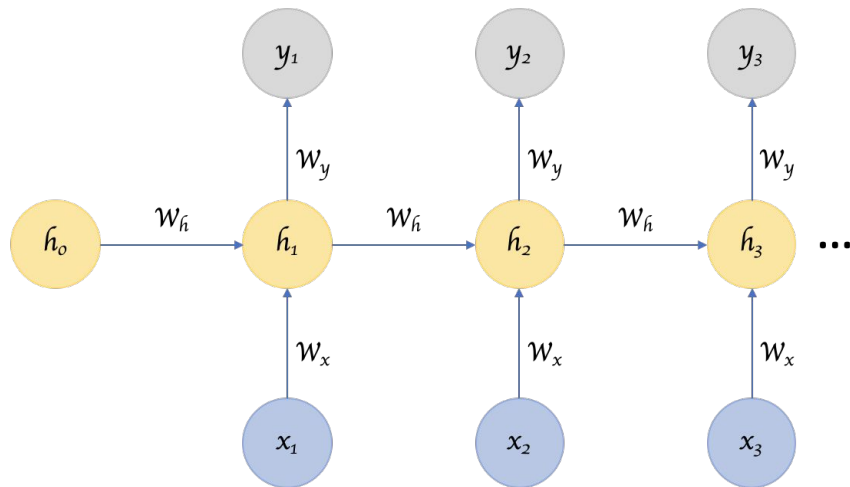
- Recurrent Neural Networks
  - Definition
  - Types of RNNs
  - Backpropagation in RNNs
- Time Series Forecasting
  - Naive Forecasting
  - MLP
  - Simple RNN
  - Deep RNNs
- Forecasting several times ahead
  - One step at a time
  - Predicting several values
  - Training for all steps
  - Confidence intervals
- Adding memory
  - LSTMs
  - GRUs
- Preview of Convolutions



LSTM  
(Long-Short Term Memory)

# Recurrent Neural Networks


# Definition




$$y_t = \mathbf{W}_y \mathbf{h}_t = \mathbf{W}_y f(\mathbf{W}_x \mathbf{x}_t + \mathbf{W}_h \mathbf{h}_{t-1})$$

# Types of RNNs


one to one




one to many




many to one




many to many



many to many



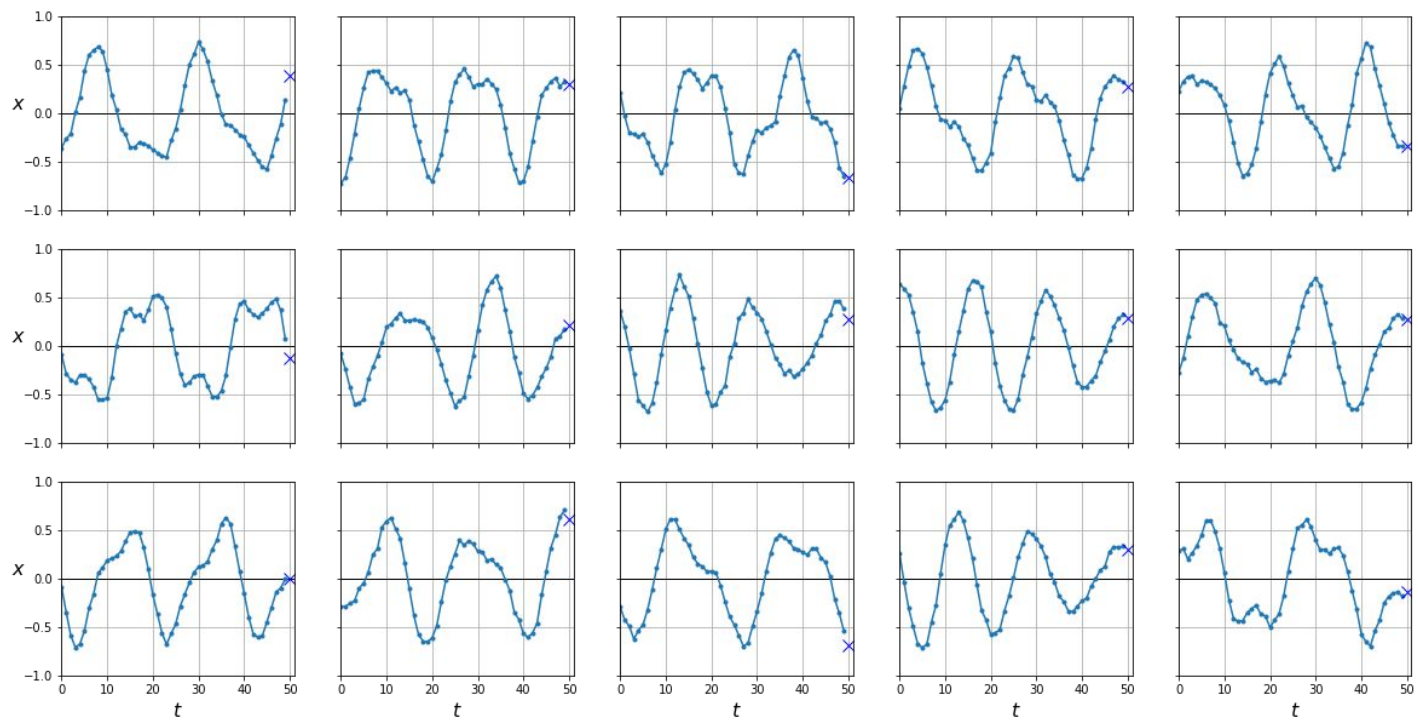
# Backpropagation in RNNs



Backpropagation Through Time



# Time Series Forecasting



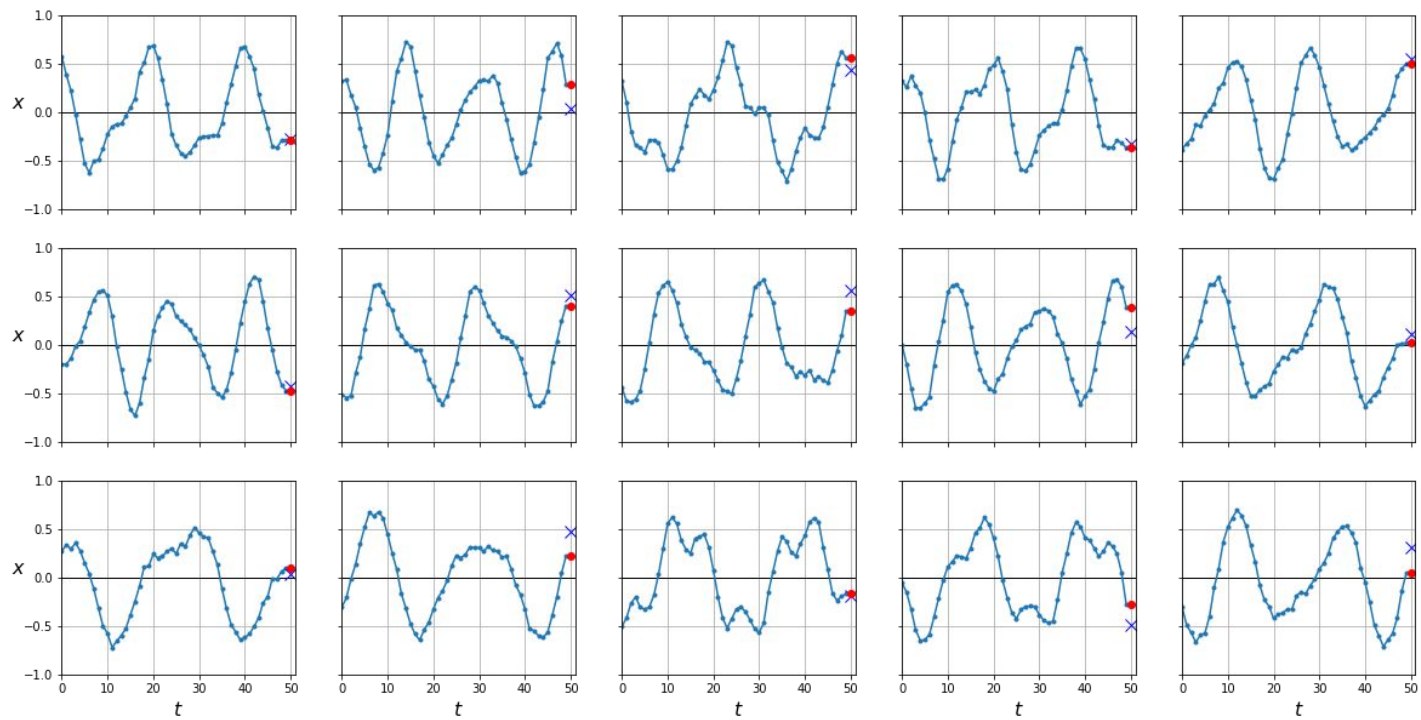


# Naive forecasting

Just predict the last value.

```
y_pred = X_test[:, -1]

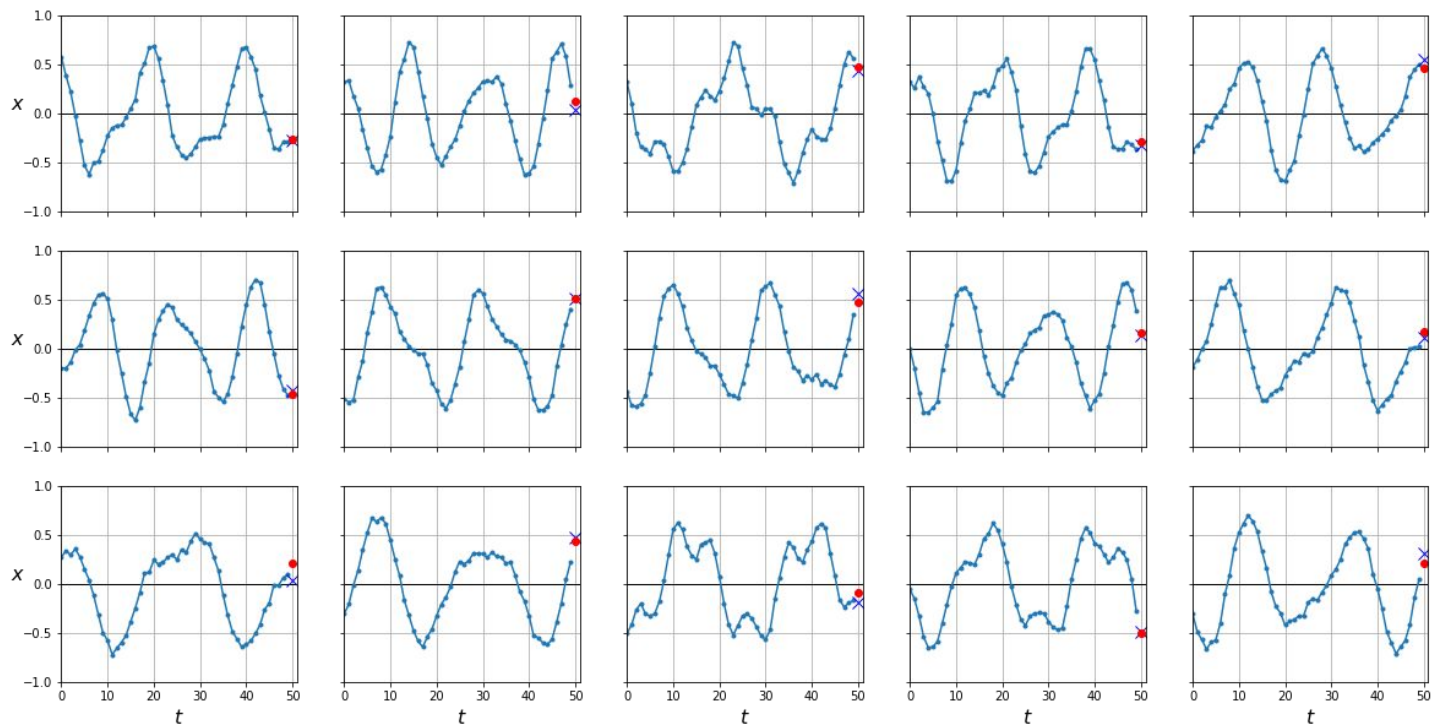
mean_squared_error(y_test, y_pred)
> 0.02157
```



# MLP

```
class MLP(torch.nn.Module):  
    def __init__(self, n_in=50, n_out=1):  
        super().__init__()  
        self.fc = torch.nn.Linear(n_in, n_out)  
  
    def forward(self, x):  
        x = x.view(x.shape[0], -1)  
        x = self.fc(x)  
        return x
```

```
mean_squared_error(y_test, y_pred)  
> 0.00481
```



# Simple RNN

```
class SimpleRNN(torch.nn.Module):  
    def __init__(self):  
        super().__init__()  
        self.rnn = torch.nn.RNN(input_size=1, hidden_size=1, num_layers=1, batch_first=True)  
  
    def forward(self, x):  
        x, h = self.rnn(x)  
        # return the last output  
        return x[:,-1]  
  
mean_squared_error(y_test, y_pred)  
> 0.02253
```





















```
class RNN(torch.nn.Module):
    def __init__(self):
        super().__init__()
        self.rnn = torch.nn.RNN(input_size=1, hidden_size=20, num_layers=1, batch_first=True)
        self.fc = torch.nn.Linear(20, 1)

    def forward(self, x):
        x, h = self.rnn(x)
        # get the last output and apply linear layer
        y = self.fc(x[:,-1])
        return y


mean_squared_error(y_test, y_pred)
> 0.00343
```



# Deep RNNs









a) 2-layer Recurrent Neural Network (RNN)



b) Unfolded 2-layer Recurrent Neural Network (RNN)





The background of the slide is a dark blue to purple gradient, overlaid with numerous out-of-focus, glowing circles in shades of light blue and cyan, creating a bokeh effect.
















**Forecasting  
several times  
ahead**



# Naive forecasting

```
y_pred = X_test[:, -1]
for step_ahead in range(9):
    y_pred = np.concatenate([y_pred, X_test[:, -1]], axis=1)

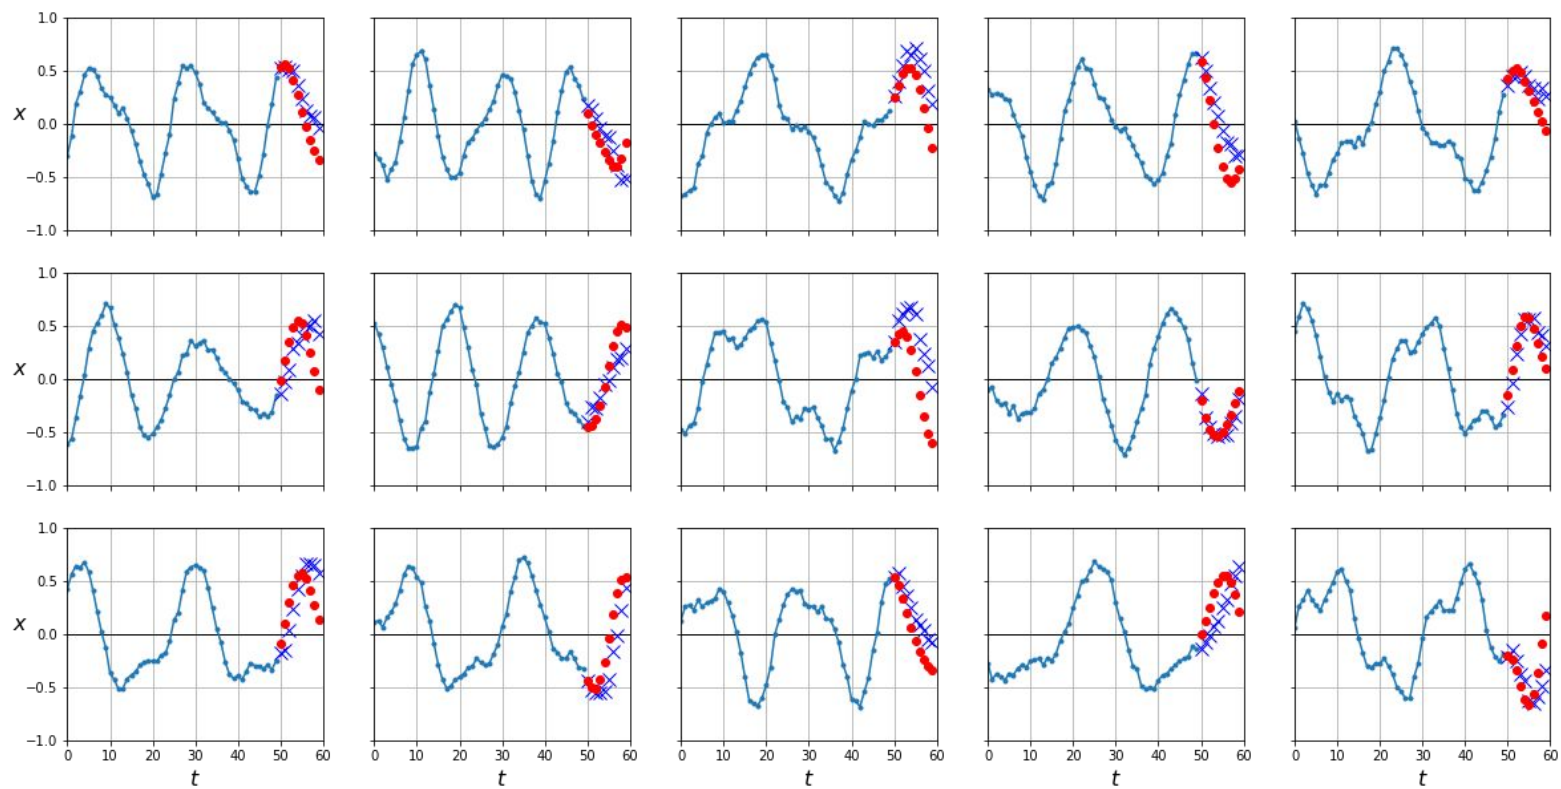
mean_squared_error(Y_test, y_pred)
> 0.2627
```



# Deep RNN (one step at a time)

```
X = X_test
for step_ahead in range(10):
    inputs = torch.from_numpy(X[:, step_ahead:]).unsqueeze(0)
    y_pred_one = model.predict(inputs).cpu().numpy()
    X = np.concatenate([X, y_pred_one[:, np.newaxis, :]], axis=1)

mean_squared_error(Y_test, y_pred)
> 0.05195
```

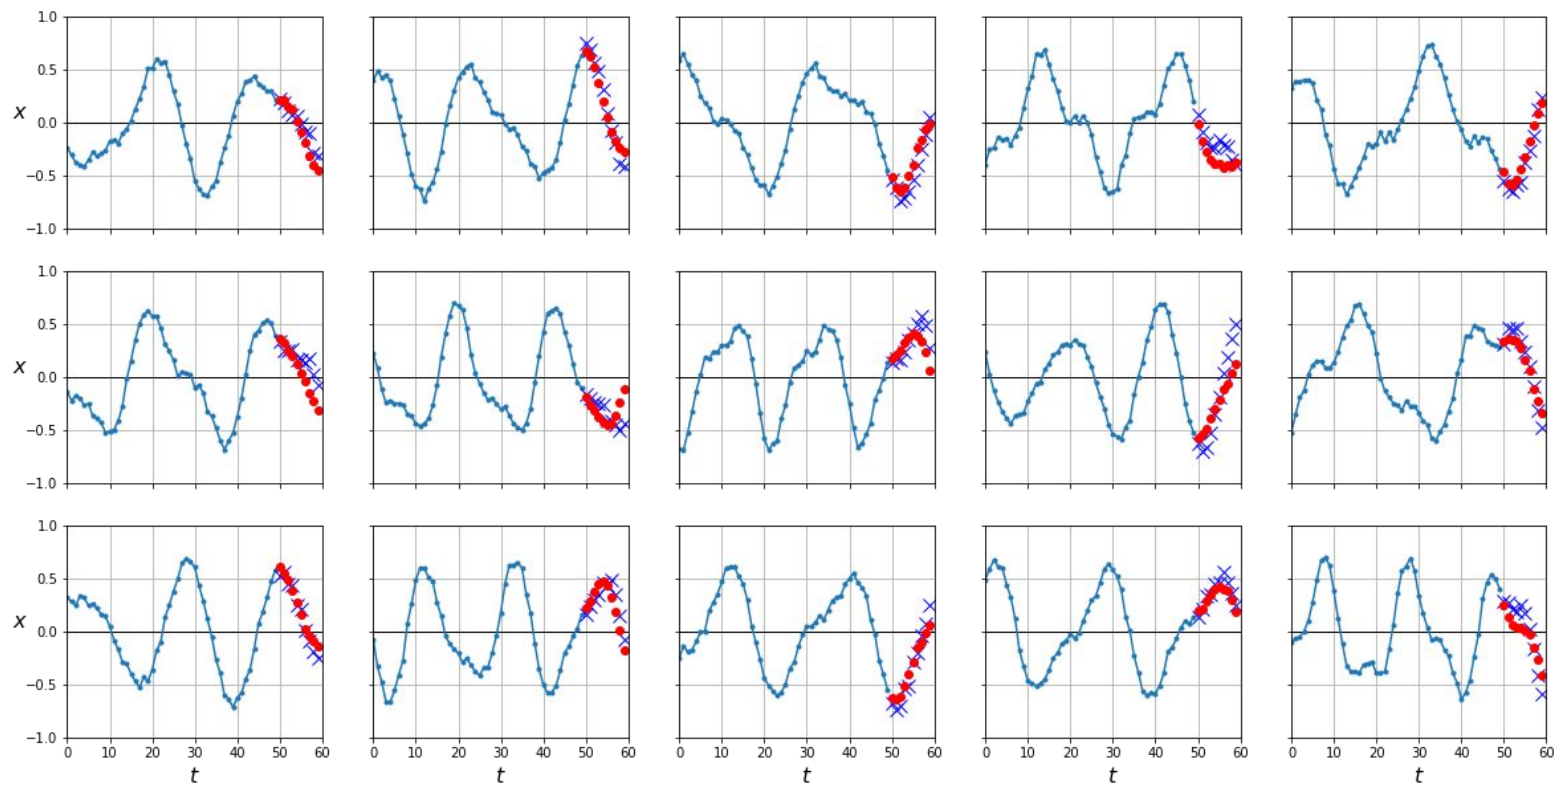


# Predicting several values

```
class DeepRNN(torch.nn.Module):
    def __init__(self, n_out=10):
        super().__init__()
        self.rnn = torch.nn.RNN(input_size=1, hidden_size=20, num_layers=2, batch_first=True)
        self.fc = torch.nn.Linear(20, n_out)

    def forward(self, x):
        x, h = self.rnn(x)
        x = self.fc(x[:,-1])
        return x

mean_squared_error(Y_test, y_pred)
> 0.011771
```



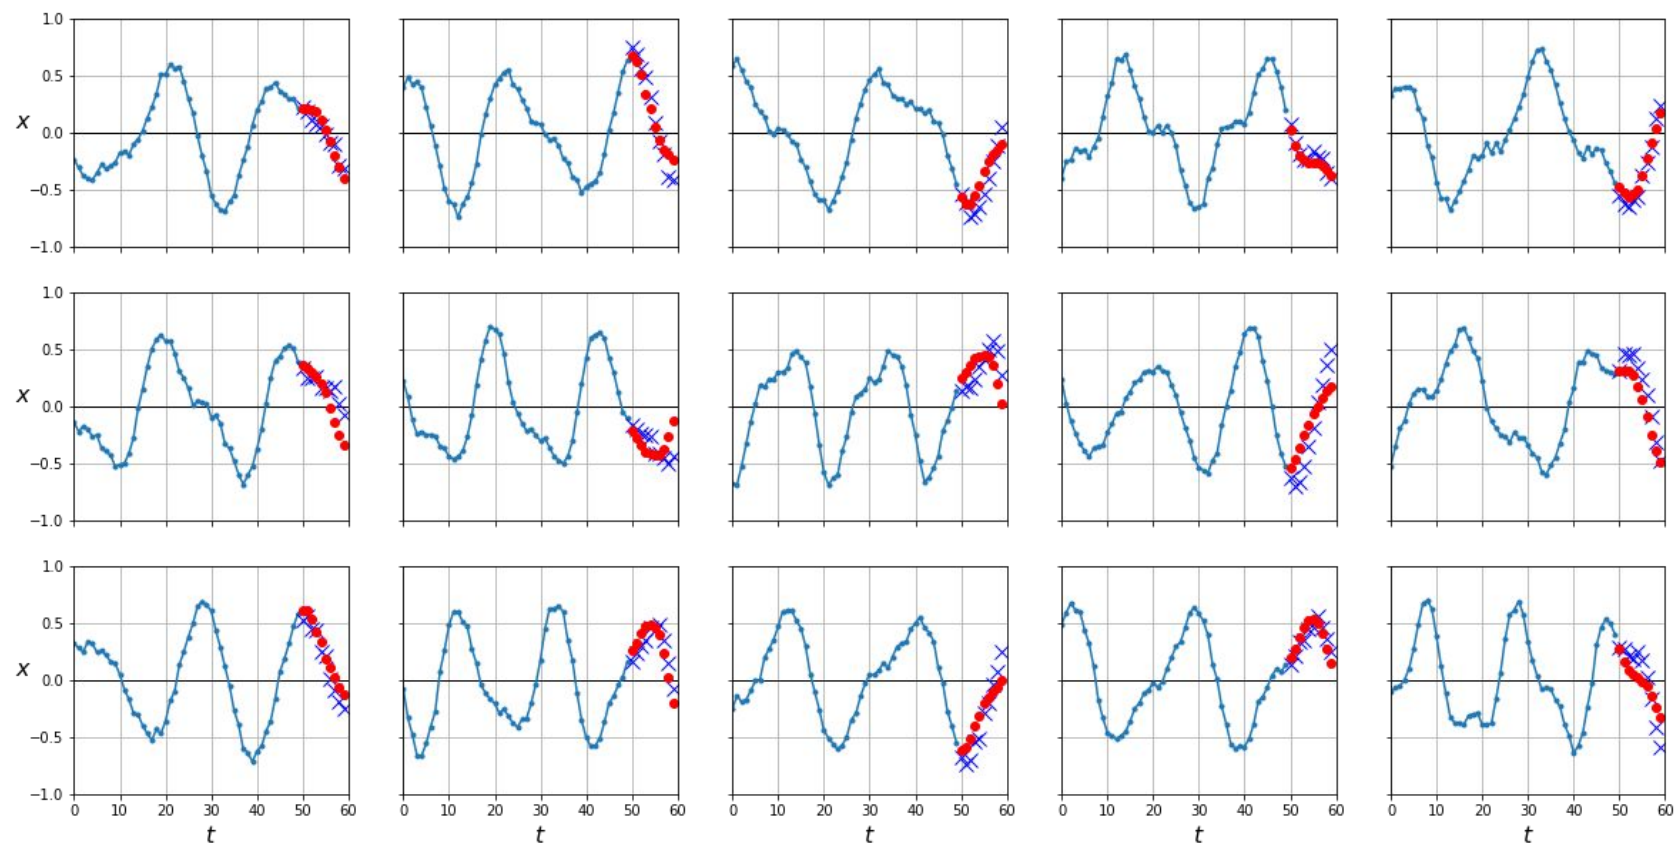


# Training for all steps

```
class DeepRNN(torch.nn.Module):
    def __init__(self, n_out=10):
        super().__init__()
        self.rnn = torch.nn.RNN(input_size=1, hidden_size=20, num_layers=2, batch_first=True)
        self.fc = torch.nn.Linear(20, n_out)

    def forward(self, x):
        x, h = self.rnn(x)
        # reshape rnn output to feed fc
        # [ Batch, time steps, features ] --> [ Batch x time steps, features ]
        x_reshaped = x.contiguous().view(-1, x.size(-1))
        y = self.fc(x_reshaped)
        # reset to original shape
        # [ Batch x time steps, features ] --> [ Batch, time steps, features ]
        y = y.contiguous().view(x.size(0), -1, y.size(-1))
        return y

mean_squared_error(Y_test[:, -1], y_pred[:, -1])
> 0.015259
```






# Confidence intervals

We can provide confidence intervals with *MC Dropout* (do predictions with dropout on and average)


```
class DeepRNN(torch.nn.Module):
    def __init__(self, n_out=10, dropout=0.5):
        super().__init__()
        self.rnn = torch.nn.RNN(input_size=1, hidden_size=20, num_layers=2, dropout=dropout,
batch first=True)
        self.fc = torch.nn.Linear(20, n_out)

    def forward(self, x):
        x, h = self.rnn(x)
        # reshape rnn output to feed fc
        # [ Batch, time steps, features ] --> [ Batch x time steps, features ]
        x_reshaped = x.contiguous().view(-1, x.size(-1))
        y = self.fc(x_reshaped)
        # reset to original shape
        # [ Batch x time steps, features ] --> [ Batch, time steps, features ]
        y = y.contiguous().view(x.size(0), -1, y.size(-1))
        return y
```



The background of the slide features a bokeh effect with numerous overlapping circles in shades of blue and purple, creating a soft, out-of-focus light pattern.

# Adding memory

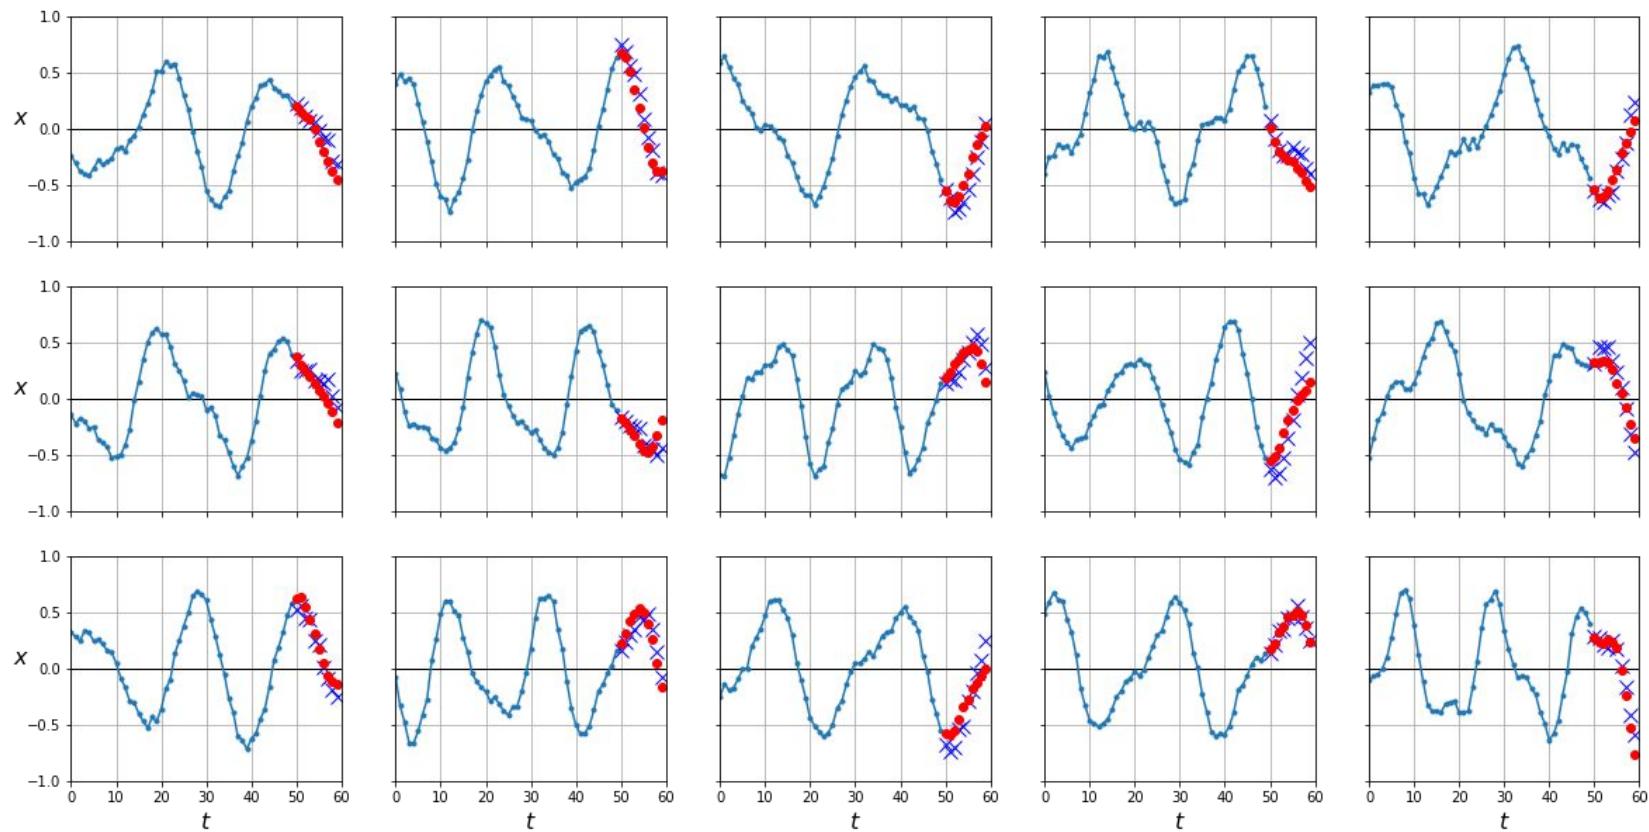



LSTM  
(Long-Short Term Memory)



# LSTM

```
class LSTM(DeepRNN):  
    def __init__(self, n_out=10, dropout=0):  
        super().__init__()  
        self.rnn = torch.nn.LSTM(input_size=1, hidden_size=20, num_layers=2,  
dropout=dropout, batch_first=True)  
  
mean_squared_error(Y_test[:, -1], y_pred[:, -1])  
> 0.00887
```








**GRU**  
**(Gated Recurrent Unit)**

# GRU

```
class GRU(DeepRNN):  
    def __init__(self, n_out=10, dropout=0):  
        super().__init__()   
        self.rnn = torch.nn.GRU(input_size=1, hidden_size=20, num_layers=2, dropout=dropout,  
batch_first=True)  
  
mean_squared_error(Y_test[:, -1], y_pred[:, -1])  
> 0.01165
```


















# Preview of Convolutions

We will talk about CNNs later, but for now keep in mind that they can also be used to process sequence data, often achieving better results.

```
class ConvRNN(torch.nn.Module):
    def __init__(self, n_out=10, dropout=0):
        super().__init__()
        self.conv = torch.nn.Conv1d(1, 20, 4, stride=2, padding=0)
        self.rnn = torch.nn.GRU(input_size=20, hidden_size=20, num_layers=2,
                                dropout=dropout, batch_first=True)
        self.fc = torch.nn.Linear(20, n_out)

    def forward(self, x):
        x = self.conv(x.permute(0,2,1))
        x, h = self.rnn(x.permute(0,2,1))
        x_reshaped = x.contiguous().view(-1, x.size(-1))
        y = self.fc(x_reshaped)
        y = y.contiguous().view(x.size(0), -1, y.size(-1))
        return y

mean_squared_error(Y_test[:, -1], y_pred[:, -1])
> 0.009064
```

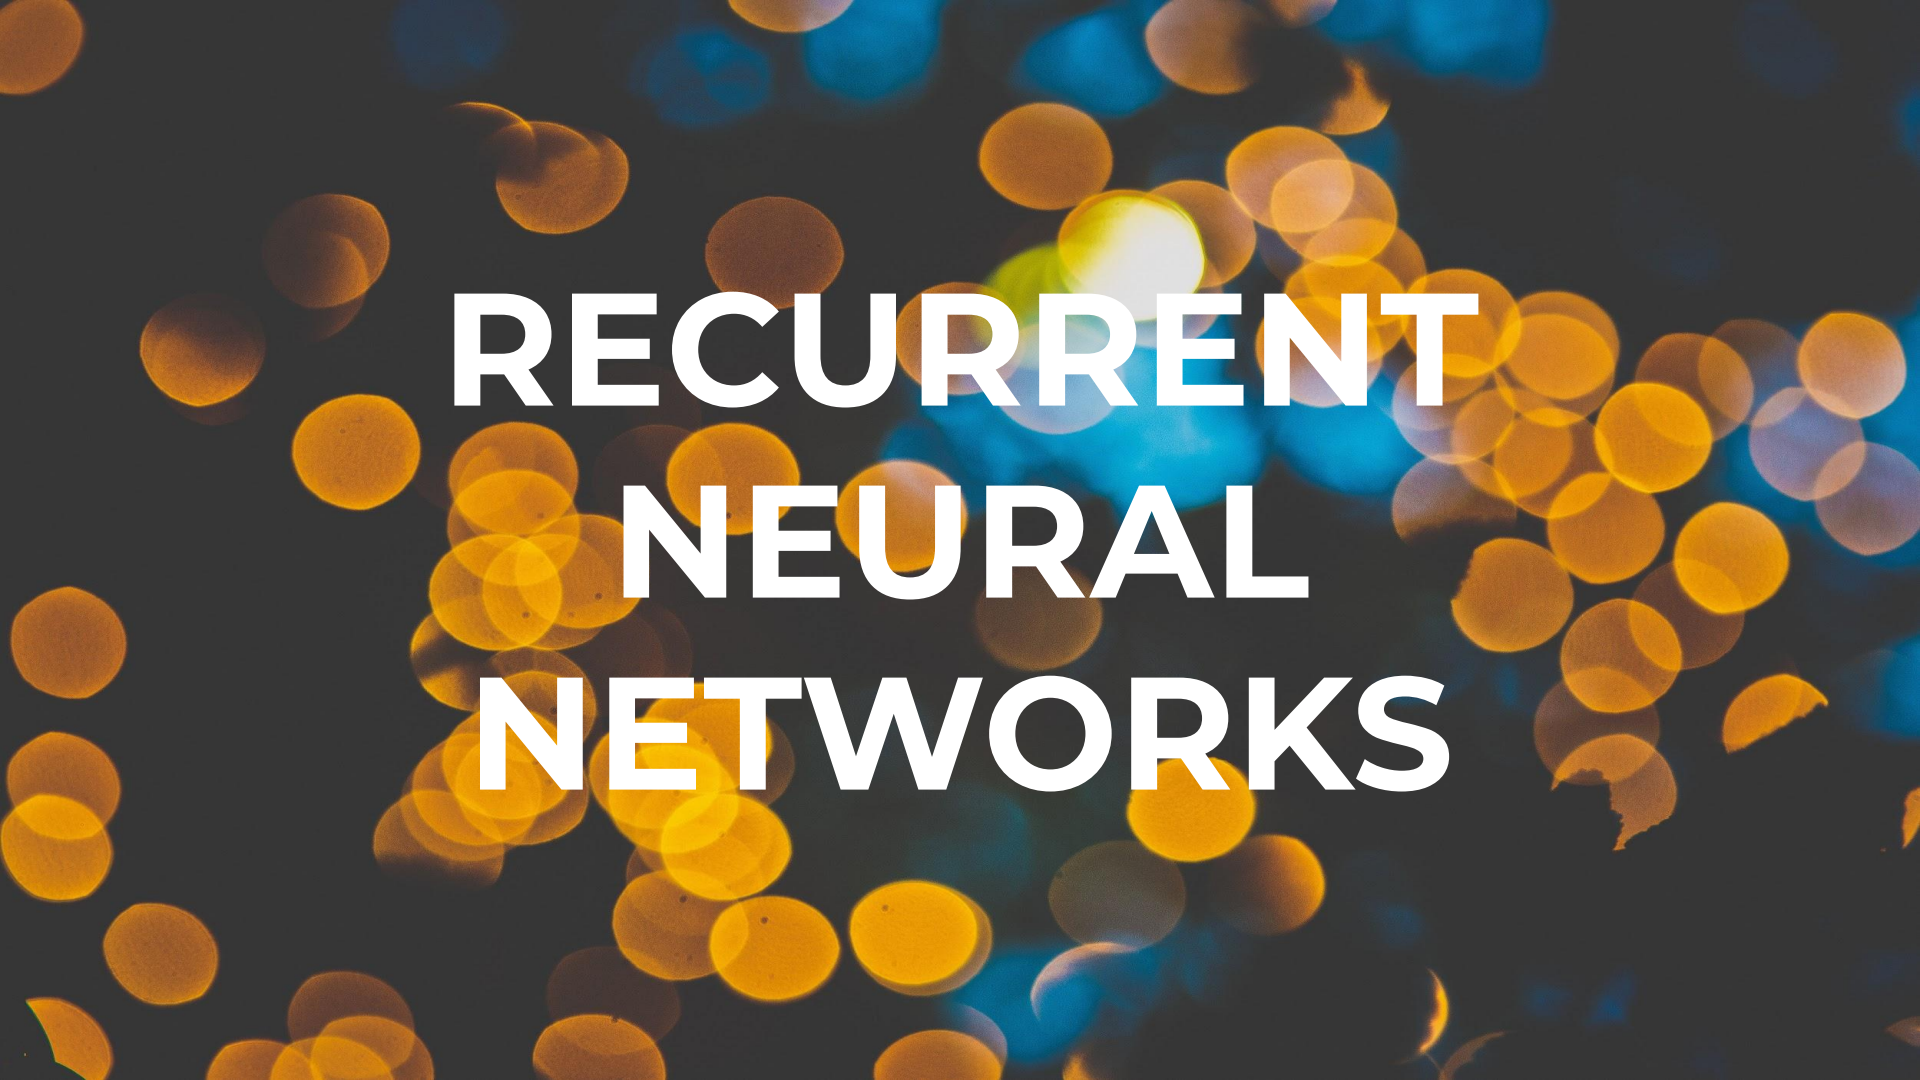




<https://github.com/sensioai/dl/blob/master/rnns/rnns.ipynb>



[https://www.youtube.com/watch?v=\\_h66BW-xNgk](https://www.youtube.com/watch?v=_h66BW-xNgk)

The background of the image is a dark, textured surface covered with numerous out-of-focus light circles, known as bokeh. These circles are primarily in shades of warm orange and yellow, with some cooler blue and teal tones interspersed, particularly towards the right side. The circles vary in size and brightness, creating a dynamic and visually appealing pattern.

# RECURRENT NEURAL NETWORKS